

The Wrong Unit of Uncertainty: Simultaneous Conformal Bands for Repeated Cross-National Surveys

The Authors

Abstract

Comparative research routinely reads a single wave-pair mean contrast and narrates it as a persistent, distribution-wide, country-level trend—attaching uncertainty to the wrong unit and over-detecting change. We give an instrument and a deployable procedure: a finite-sample simultaneous band over a country’s whole attitude *trajectory*, with the country as the exchangeable unit and a hierarchy of claims whose top rung is the honest object. The band is a clustered conformal band exact at any K , inside a provably nominal-safe selector that corrects for survey design only when it can be certified reliable and otherwise abstains. Applied to real data it changes the substantive picture: in the ESS (2018–2022), persistent distribution-wide trust decline is certified in one of roughly thirty democracies (Greece), not the twenty flagged marginally; on the World Values Survey (1981–2022) the marginal case for Foa–Mounk “deconsolidation” over-counts the persistent phenomenon by $2.6\text{--}6.5\times$, and the deconsolidation that survives concentrates in post-communist and Arab-Spring-aftermath states rather than the consolidated West the thesis targeted. A second, subtler error—the calibration objects, other countries’ distributions, are estimated from complex samples—we resolve by mapping its scope: correcting for it is non-identified without the design-noise law and, even with it, unreachable at survey scale, so the abstaining reduction is the right

object rather than a concession.

1 Introduction

Comparative politics increasingly treats a national public as a *distribution* rather than a mean. Trust in parliament, satisfaction with democracy, or support for the political system in a country and year is a whole cumulative distribution function over a response scale, and repeated cross-national surveys—the European Social Survey (ESS; Kaminska and Lynn, 2017), the World Values Survey, the AmericasBarometer (Cohen et al., 2021)—let us track how that distribution moves across waves and, increasingly, *transport* it: predict one country’s trajectory from the others, or flag which countries depart from a common pattern. Prediction bands and conformal methods are a natural tool here, promising finite-sample, distribution-free coverage.

Two wrong units. A claim and its calibration can each be attached to the wrong object. The *claim* “democratic trust is declining” is typically built from a single wave-pair mean contrast, then narrated as a persistent, distribution-wide, country-level trend. Requiring the whole distribution to shift, *simultaneously* across the entire trajectory, is a strictly stronger and more honest object, and the number of countries that qualify collapses between the two. Our instrument is a finite-sample simultaneous band over a country’s whole trajectory, with the country as the exchangeable unit, and an ordered hierarchy of claims (pairwise → any-pair → net → persistent country-wide → Bonferroni-across-countries) whose top rung is the object casual readings implicitly assert. This multiplicity-honest core needs no survey design at all and applies to any repeated distributional claim.

The second, subtler wrong unit is a layer of uncertainty these exercises silently omit. Conformal inference is usually calibrated on *observed outcomes*. Repeated cross-national surveys are different: the calibration objects—the other countries’ attitude distributions—are themselves *not observed*. Each is estimated from a complex probability sample with

weights, clustering (primary sampling units), and stratification, and therefore carries its own sampling error. Standard practice plugs these survey estimates into the calibration step as if they were the truth, so the transport band, honest about variation *across* countries, stays blind to the fact that each “country” is a noisy estimate of a latent population function.

The two errors are genuinely distinct. Writing $F_{c,r}(t)$ for country c ’s latent attitude CDF at threshold t in wave r , and $\tilde{F}_{c,r}(t)$ for its complex-survey estimate,

$$\tilde{F}_{c,r}(t) - \mu_r(t) = \underbrace{(F_{c,r}(t) - \mu_r(t))}_{\text{transport error}} + \underbrace{(\tilde{F}_{c,r}(t) - F_{c,r}(t))}_{\text{design error}}, \quad (1)$$

where μ_r is the transport center. A conformal band calibrated on the observed $\tilde{F}$ scores treats the sum as if it were the transport error alone: when the design error is non-negligible the band is systematically too wide, yet any correction that subtracts the survey noise can make it too narrow, because the correction is itself estimated from a small number of countries.

Contributions. We make four. *First, a diagnosis and a portable instrument.* We separate the two wrong units—the claim unit (a wave-pair contrast standing in for a persistent distribution-wide trajectory) and the calibration unit (estimated distributions treated as latent, the omitted design-error layer of (1))—and give a finite-sample simultaneous band over a country’s whole trajectory, with a multiplicity-honest guarantee hierarchy that corrects the first. *Second, a provably-valid deployed procedure.* The band we recommend is a clustered “population” conformal band (PCB), finite-sample exact at *any* K (Theorem 3), wrapped in a nominal-safe selector that corrects for the design error only when four target-blind gates certify reliability and otherwise reduces to the exact band or abstains conservatively (Theorem 5). We confirm the frozen procedure on an independent, pre-registered simulation spanning ten data-generating processes it has never seen. *Third, two named re-analyses* (Section 8) where the instrument changes the substantive picture. For European trust (ESS 2018–2022), marginal wave-by-wave estimates suggest widespread decline, but a

country-wide, design-aware simultaneous band certifies a *persistent distribution-wide* decline in only one country (Greece), not the twenty flagged marginally. Globally, we reanalyze the Foa–Mounk “democratic deconsolidation” thesis on the World Values Survey (1981–2022) and find its marginal evidence over-counts persistent, distribution-wide deconsolidation by $2.6\text{--}6.5\times$ —real in a specific minority, not the broad erosion the marginal reading implies. Aggregating the per-item certifications turns the correction into a discovery: the *certified core* of multi-item deconsolidation (thirteen countries) is concentrated in post-communist and Arab-Spring-aftermath states, with the consolidated West entering only through rising openness to strongman and army rule (Finland, Switzerland)—the deconsolidation that survives the honest bar lives largely outside the mature democracies the thesis was formulated about. *Fourth, the scope condition that justifies this deployment.* We delineate exactly when a design-aware correction is available and when it is not, so that abstention is a principled choice rather than a concession. Without a model of the design-noise law the latent transport-error distribution is non-identified—a deconvolution problem (Theorem 1); and even *with* that information the correction is *unreachable* at cross-national scale, for two quantitative reasons: the design-to-total noise ratio ρ saturates well below the level at which deconvolution helps, and the deployed reliability diagnostic obeys the algorithmic floor $D \geq \sqrt{2/(K-1)}$ —an identity of the frozen procedure, holding for every possible dataset—that needs $K \geq 94$ exchangeable populations. On the two largest such surveys these barriers bind in *opposite* ways—the World Values Survey clears the size barrier ($K \approx 100$) but carries only weights-based, negligible design noise, while the ESS has genuine clustered design noise but far too few countries—so neither clears both. This is why the abstaining reduction above, not a design-aware correction, is the right object at survey scale. The empirical payoff is thus the two reanalyses; the design-aware layer is a diagnostic and scope result that, on real survey data, certifies the simple band rather than firing—so the deconvolution is not something the reader should expect to activate empirically.

The remainder situates the contribution among the conformal and political-trust litera-

tures (Section 2), sets up the problem and impossibility result (Section 3), builds the safe adaptive method and its theory (Sections 4–5), validates it in simulation (Section 6), characterizes the real-data regime (Section 7), reanalyzes political trust (Section 8), and discusses scope (Section 9).

2 Related work

Our work sits at the intersection of four literatures. First, **distribution-free conformal prediction** gives finite-sample coverage without modeling assumptions (Vovk et al., 2005; Lei et al., 2018). Because political data rarely satisfy exchangeability, we build on its relaxations: weighted conformal prediction for covariate shift (Tibshirani et al., 2019) and nonexchangeable conformal prediction, which bounds the coverage gap by the total variation between exchanged distributions (Barber et al., 2023). Gibbs and Candès (2021); Gibbs et al. (2024) adapt conformal sets to distribution shift and interpolate between marginal and conditional validity. Our “adaptive” selector differs in kind: it adapts not to temporal shift but to the *certified magnitude of survey design noise*, reducing to a fixed clustered method when that noise is negligible. Within the discipline, Randahl et al. (2026) bring conformal prediction to *Political Analysis* for the first time, calibrating bin-conditional intervals around a point forecast of conflict fatalities; our object is different in kind—finite-sample *simultaneous* trajectory bands with the country as the exchangeable unit and estimated calibration distributions, rather than bin-conditional intervals on a single prediction.

Second, we treat the **population (country) as the exchangeable unit**, following hierarchical conformal prediction for two-layer models (Dunn et al., 2023) and distribution-free inference with hierarchical data (Lee et al., 2023). The same finite-sample-exact clustered construction is developed, for a different (non-survey) domain, in a companion paper by the present authors (reference omitted for anonymous review). These methods transport guarantees to an unseen cluster—our cross-national setting—but treat each cluster’s data

as cleanly sampled. Our departure is that the calibration objects are not observed: each country’s attitude distribution is *estimated* under a complex survey design, so design-based sampling variance (Binder, 1983; Lumley, 2010) contaminates the very scores conformal calibration relies upon.

Third, the closest methodological neighbors handle **noisy or estimated calibration**. Wieczorek (2023) develops design-based conformal prediction under complex sampling, but design enters through test-point weighting rather than as noise in transported objects. Einbinder et al. (2024) and Sesia et al. (2025) correct conformal calibration under label noise, the latter by deconvolving a *known* noise model. Our central theoretical result is that this identification premise generally fails for survey design noise: recovering the latent transport-error law is a deconvolution problem (Fan, 1991) that is non-identified without a design-noise law, so no observed-score band is simultaneously valid and strictly narrower than the conservative plug-in. And even when a noise model *is* supplied, we show the deconvolution’s reliability is itself bounded at a finite number of calibration units by an algorithmic floor (Proposition 1)—a finite- K obstruction distinct from the asymptotic identifiability the classical deconvolution literature studies.

Fourth, our target is a **functional object**—a CDF over thresholds tracked across rounds—placing us near simultaneous functional conformal bands (Diquigiovanni et al., 2021). This is where we part from generic multiplicity control. Procedures like FDR or Bonferroni (Benjamini and Hochberg, 1995) adjust a collection of tests after the fact; our guarantee is instead carried by the geometry of a single finite-sample band—a supremum over wave-pairs and thresholds with the country as the exchangeable unit—so the “persistent, distribution-wide” object is *defined and covered directly*, not recovered by correcting marginal p -values. The distinction is one of estimand, not merely error rate, and partial-pooling cautions about reading marginal country trends (Gelman et al., 2012) motivate but do not deliver the stronger claim. Substantively, our two reanalyses supply the formal instrument a largely qualitative skepticism has lacked, engaging debates over whether European political trust is in broad de-

cine (Norris, 2017; Voeten, 2017; Wuttke et al., 2022; Valgarðsson et al., 2025) and whether publics are “deconsolidating” (Foa and Mounk, 2016, 2017).

3 Setup and the two targets

The object we deploy is the exact simultaneous band of Theorem 3; the impossibility below (Theorem 1) is a scope boundary delimiting when a design-aware correction to that band is even available.

Objects. We observe K source countries, indexed $c = 1, \dots, K$, treated as exchangeable units (the “population” or cluster is the exchangeable object, not the individual respondent). Each country has a latent attitude trajectory $F_{c,r}(t)$, a CDF over a fixed grid of thresholds $t \in \{t_1, \dots, t_T\}$ and survey rounds r . We do not see $F_{c,r}$; we see a complex-survey estimate $\tilde{F}_{c,r}(t) = F_{c,r}(t) + S_{c,r}(t)$, where $S_{c,r}$ is design (sampling) error with a design variance $v_{c,r}(t)^2$ that a stratified-PSU / replicate-weight bootstrap can estimate. Define the transport score of a country as its worst-case standardized deviation from the transport center,

$$R_c = \max_{r,t} \frac{|F_{c,r}(t) - \mu_r(t)|}{\sigma(t)}, \quad \tilde{R}_c = R_c + \xi_c,$$

where $\tilde{R}_c$ is the same quantity computed from the observed $\tilde{F}$ and ξ_c is the induced design contamination. The scalar that governs the whole problem is the design-to-*total* standard-deviation ratio

$$\rho = \frac{\text{SD}(\xi)}{\text{SD}(\tilde{R})} = \sqrt{\frac{\overline{v^2}}{s_R^2 + \overline{v^2}}} \in [0, 1),$$

with $\rho \rightarrow 0$ the design-noise-negligible regime and $\rho \rightarrow 1$ the regime in which survey noise dominates the between-country spread. (The generative dial in the simulations is the unbounded design/transport ratio $\text{SD}(\xi)/\text{SD}(R)$; all reported ρ values, gates, and the $\sqrt{1 - \rho^2}$ width law use the bounded design/total definition above.)

Target and validity. We want a simultaneous band $\widehat{B}(r, t)$ that covers a *new* country’s *latent* trajectory, $\Pr(F_{\text{new}, r}(t) \in \widehat{B}(r, t) \forall r, t) \geq 1 - \alpha$. Two facts frame everything. Ordinary clustered conformal on the observed scores is exact for the *survey-estimate* target $\widetilde{F}_{\text{new}}$ (it only needs exchangeability of countries); its coverage of the *latent* target differs by an explicit gap that vanishes with the design-noise scale (Theorem 2(b)—symmetry of the design law alone is *not* enough for conservativeness, so we bound the gap rather than assume its sign). Meanwhile its width is inflated by the design noise, by a factor that grows with ρ . Correcting the width requires estimating and removing the design contribution—a deconvolution.

Country exchangeability. Exchangeability of countries is our sole identifying assumption, and it is weaker than the i.i.d.-respondents assumption standard practice already makes within each survey. When it holds only approximately, coverage degrades gracefully: the nonexchangeable conformal bound of Barber et al. (2023) controls the coverage gap by the total variation between the exchanged country distributions, so moderate departures cost a small, bounded amount rather than invalidating the band. The conservative (Bonferroni-across-countries) rung of the guarantee hierarchy further protects the “only Greece” claim against the family-wise multiplicity of screening many countries at once—a distinct safeguard from the exchangeability bound above, which controls non-exchangeability directly and which Bonferroni does not replace.

The impossibility. The central obstruction is that the deconvolution is not identified from the contaminated scores alone.

Theorem 1 (Non-identification and necessity of design information). *Let the observed score be $\widetilde{R} = R + \xi$ with $R \perp \xi$. (i) The law of R is not identified from the law of $\widetilde{R}$ without restrictions on the law of ξ : there exist distinct pairs $(\mathcal{L}_R, \mathcal{L}_\xi) \neq (\mathcal{L}'_R, \mathcal{L}'_\xi)$ with $\mathcal{L}_R * \mathcal{L}_\xi = \mathcal{L}_{R'} * \mathcal{L}_{\xi'}$. (ii) Consequently, among bands that are measurable functions of the observed scores only, none is valid for the latent target at level $1 - \alpha$ for all admissible design laws while being strictly narrower than the plug-in band whose radius is the $(1 - \alpha)$ conformal*

quantile of $|\tilde{R}|$. In the degenerate case $\xi \equiv 0$ the plug-in band is itself the tightest valid band, so no uniform improvement exists.

The proof (Appendix A, with a numerical contract test) is a convolution-invertibility argument for (i) and a conformal converse for (ii). The reading is constructive: a band both honest and narrower than the plug-in must *supply* the design-noise law, which complex surveys do provide—through replicate weights or a design bootstrap, exactly what Section 4 uses—so this input is not optional.

Remark 1 (Beyond surveys). Nothing in Theorem 1 is specific to surveys: it holds whenever a conformal procedure is calibrated on *estimated* objects, so that the nonconformity score is a noisy version $\tilde{R} = R + \xi$ of a latent score—model-based small-area (MRP) outputs, learned or imputed features, meta-analytic effects. The same two-part consequence follows: design/measurement information is the identifying ingredient, and its reliability at a finite number of calibration units is itself bounded (Proposition 1). The barrier is confirmed outside surveys on a model-based small-area calibration (Supplementary Material).

4 A nominal-safe adaptive method

Theorem 1 says the design bootstrap is the identifying input. The method turns that input into a band by choosing, target-blind, among three constructions and by never using the aggressive one unless it is provably reliable at the available number of countries.

Three branches. For calibration error curves $\{\tilde{E}_c\}_{c=1}^K$ over threshold grid t , with per-country design SD $\hat{v}_c(t)$ from the stratified-PSU bootstrap: (i) *Clustered PCB* takes the conformal quantile of the *unstudentized* scores $\max_t |\tilde{E}_c(t)|$ (the U0 construction of Theorem 3); it is finite-sample exact for the survey-estimate target at any K , and its latent-target gap is the explicit $O(\rho)$ bound of Theorem 2(b). (ii) *Deconvolution* studentizes by the finite- K -safe target scale $\hat{s}_T^2(t) = \max(s_{\text{plug}}^2(t) - [\widehat{v}^2(t) - z \text{ SE}]_+, \text{floor})$, which subtracts a *lower*

confidence bound on the design variance so the correction cannot over-shrink; its width scales as $\sqrt{1 - \rho^2}$ relative to PCB. (iii) *Conservative* inflates each unstudentized score by the design SD, $\max_t(|\widehat{E}_c(t)| + z\widehat{v}_c(t))$ — a score that dominates the PCB score pointwise, so the conservative band *contains* the PCB band by construction. This nesting is load-bearing: selection between nested bands anchored at an exact band is validity-free (Theorem 5), so the selector needs no conditional-coverage argument.

The three questions. The naive selector asks only *can we deconvolve* and *do we need to*. The finite- K failure of deconvolution (Section 6) comes from a missing third question: *can we do it safely at this K ?* The safe selector adds it explicitly.

Safe selector (Algorithm 1). Let $\widehat{\rho}_{\text{LCB}}$ be a conservative lower bound on ρ , $D = \max_t \text{SE}(\widehat{s}_T^2)/\widehat{s}_T^2$ a reliability diagnostic, and $\widehat{\delta}_{\text{UCB}}(D)$ a frozen upper bound on the finite- K coverage remainder. Deconvolution is used *iff all four gates pass*: (A) $\widehat{\rho}_{\text{LCB}} > \rho_0$ (design noise materially large); (B) $\widehat{\delta}_{\text{UCB}}(D) \leq \delta_{\text{max}}$ (remainder within budget); (C) $\widehat{\text{Gain}} - \Delta_g > g_{\text{min}}$ (beats conservative by a margin); (D) stability. Otherwise the band is PCB when $\widehat{\rho}_{\text{LCB}} \leq \rho_0$ and the conservative envelope when design noise is real but information insufficient. Every gate is a function of the calibration set only, so the selector is target-blind — though, by the nested-band argument of Theorem 5, validity would survive even an adversarial selection between the two anchor branches. The deconvolution branch, which is *not* nested, is charged its own miss budget $\alpha_{\text{dec}} = \max(0.1\alpha, 3/(K+1))$ (the anchors run at $\alpha - \alpha_{\text{dec}}$), and only in the $K \geq 94$ regime where its gate can open at all; at cross-national-survey K the anchors keep the full α and the deployed guarantee is exact. All constants are frozen before validation (Section 5, Appendix B); when deconvolution is used the reported guarantee is the observable $1 - \alpha - \widehat{\delta}_{\text{UCB}}(D)$.

Deployment. The method ships as one call, `dapcb(cal_errors, v_cal, center, alpha)`, in an installable Python package (module `pcb`) provided in the replication package and to be

Algorithm 1: Nominal-safe adaptive design-aware band

Input: calibration curves $\{\tilde{R}_c\}$ (unit = country), design replicates $\{\hat{v}_c\}$, target curve $\hat{F}_{\text{new}}$, level α ; frozen $\rho_0, \delta_{\max}, g_{\min}, \Delta_g$
compute $s, \hat{s}_T$, and the three radii $r_{\text{PCB}}, r_{\text{dec}}, r_{\text{con}}$; diagnostics $\hat{\rho}_{\text{LCB}}, D, \hat{\delta}_{\text{UCB}}(D)$;
if $\hat{\rho}_{\text{LCB}} \leq \rho_0$ **then**
 | branch $\leftarrow$ PCB;
else if $\hat{\delta}_{\text{UCB}}(D) \leq \delta_{\max}$ **and** *stable* **and** $\widehat{Gain} - \Delta_g > g_{\min}$ **then**
 | branch $\leftarrow$ deconvolution;
else
 | branch $\leftarrow$ conservative;
return $\hat{F}_{\text{new}} \pm r_{\text{branch}}$, branch, and diagnostics;

archived with a version tag on the Dataverse, returning the band together with the selected branch, $\hat{\rho}, \hat{\rho}_{\text{LCB}}, D, \hat{\delta}_{\text{UCB}}$, the coverage level, and the fallback reason—so a user sees *why* the band was built as it was and reads off the finite- K guarantee it carries.

5 Theory

We state validity for the oracle, for the estimated-law procedure, and for the deployed safe selector, then the efficiency it buys. Full proofs are in Appendix A; each theorem has a matching machine-checked contract test in the replication package.

Theorem 2 (Oracle validity). *Suppose the countries are exchangeable and the design law is known. (a) The clustered band is finite-sample valid for the survey-estimate target: $\Pr(\tilde{F}_{\text{new}} \in \hat{B}) \geq 1 - \alpha$, using exchangeability only. (b) For the latent target the same band satisfies, for every $u > 0$, $\Pr(F_{\text{new}} \in \hat{B}) \geq 1 - \alpha - \Pr(\xi < -u) - \Pr(\tilde{R} \in (\hat{q} - u, \hat{q}])$: the gap between the two targets is a noise-tail term plus a local anti-concentration term, both vanishing with the design-noise scale ($O(\rho)$ under a bounded score density near the quantile), and zero by inspection when the design noise is zero. Symmetric design noise alone does not make the band conservative for the latent target — a bimodal latent score law furnishes a counterexample (Supplementary Material, Remark on the latent-target bridge) — which is why we report the observed ρ that makes the gap negligible rather than assuming conservativeness. (c) The*

deconvolution band with the true target scale is valid for the latent target with width a factor $\sqrt{1 - \rho^2} + o(1)$ of the clustered band.

Theorem 3 (Exact finite- K validity of the deployed base band). *Let the base band use the unstudentized cluster score $R_c = \max_t |E_c(t)|$ (no data-dependent modulation), with a center satisfying the symmetric-construction condition: fixed, split-fold, per-unit from that unit’s own data (the deployed own-previous-round centers), or a symmetric function of all $K+1$ curves (Appendix A states the four cases). If the countries are exchangeable, then for every K the band $\widehat{B} = \text{center} \pm \widehat{q}$, with $\widehat{q}$ the $\lceil (1 - \alpha)(K+1) \rceil$ -th order statistic of $\{R_c\}$, satisfies $\Pr(F_{\text{new}} \in \widehat{B}) \geq \lceil (1 - \alpha)(K+1) \rceil / (K+1) \geq 1 - \alpha$, distribution-free. Two asymmetries are excluded by construction and quantified when present: dividing R_c by an in-sample modulation $\widehat{s}(t) = \text{SD}_c E_c(t)$ —which the unobserved target does not share—induces an $O(g/K)$ self-inclusion deficit, and centering calibration scores on a grand mean that contains their own unit while the unsurveyed target sits outside induces an anti-conservative $O(1/K)$ bias, removed by leave-one-out centers (Appendix A). We therefore deploy the unstudentized band with symmetric centers; on the near-homoskedastic real curves the width cost of dropping the modulation is $\leq 5\%$.*

Theorem 4 (Estimated-law validity). *Assume the observed and latent score curves sit in a common scale family indexed by the design variance (A1), the sup-score law satisfies an anti-concentration bound near its upper quantile (A2), and fourth moments and an unbiased size- B design bootstrap exist (A3) — Appendix A states all three precisely. Then (i) with oracle scales the deconvolution band is finite-sample exact for the latent target at every K , and (ii) with estimated scales it satisfies $\Pr(F_{\text{new}} \in \widehat{B}) \geq 1 - \alpha - \varepsilon_{K,B}$ with $\varepsilon_{K,B} = O(\sqrt{\log(KT)/\min(K, B)})$, proved by a sandwich and concentration argument with explicit constants. The finite- K -safe scale $\widehat{s}_T^2$ (Section 4) subtracts a lower confidence bound on the design variance, making the scale-error term one-sided (wide, not narrow) outside an α_2 -probability event. Assumption (A1) is not removable: the bimodal counterexample of Theorem 2(b) defeats any latent-target band without a shape restriction of this type.*

Theorem 5 (Safe-adaptive finite- K validity—the deployed pipeline). *Deploy Algorithm 1 with unstudentized PCB and conservative branches (scores $\max_t |E_c(t)|$ and $\max_t (|E_c(t)| + z\widehat{v}_c(t))$, both at level α_{anchor}) and the deconvolution branch at its own budget α_{dec} , where $(\alpha_{anchor}, \alpha_{dec}) = (\alpha, 0)$ when the reliability gate is infeasible at the observed K (Proposition 1: $K < 94$, i.e. every cross-national survey) and otherwise $\alpha_{dec} = \max(0.1\alpha, 3/(K+1))$ with $\alpha_{anchor} = \alpha - \alpha_{dec}$, a frozen function of K alone. Then, finite-sample and with no conditions on the selection rule,*

$$\Pr(F_{new,r}(t) \in \widehat{B}(r, t) \forall r, t) \geq 1 - \alpha - \varepsilon_{K,B} \mathbf{1}\{K \geq 94\},$$

where $\varepsilon_{K,B}$ is Theorem 4’s estimated-law remainder, attached to the deconvolution branch only and reported through the observable diagnostic $\widehat{\delta}_{UCB}(D)$ when that branch fires.

The proof (Appendix A) uses no conditional-on-selection step. Its two ingredients are a nested-band lemma—the conservative score dominates the PCB score pointwise, so the two bands are nested, and *any* selection between nested bands anchored at an exact band is validity-free—and a union bound that charges the non-nested deconvolution branch its own α_{dec} miss budget, but only in the $K \geq 94$ regime where Proposition 1 shows its gate can open at all. The gates of Algorithm 1 therefore carry no validity burden; they are an efficiency device choosing among individually guaranteed constructions. At survey scale the deployed guarantee is exact $1 - \alpha$ with no remainder — the guarantee a user actually deploys — and the only remainder anywhere in the pipeline is confined to the many-unit regime where deconvolution can activate.

Efficiency. Three statements bound the price and payoff (proved in the supplement). (AW-1) in the low- ρ regime the deconvolution/PCB width ratio is $1 - \frac{1}{2}\rho^2 + O(\rho^4)$, matching the observed reduction to the decimal; (AW-2) the conservative branch is never wider than the Bonferroni envelope and at least as wide as PCB — under the deployed unstudentized construction the dominance over PCB holds pointwise by construction, making the two

Table 1: Whole-trajectory coverage (4000 reps, nominal 90%) as the trajectory length L grows. A band calibrated per point or per round covers each of its own units at the nominal rate, but the probability that it covers the *entire* trajectory decays; the country-trajectory band—the correct unit—holds nominal at every L .

L (rounds)	marginal (per point)	per-round	trajectory (ours)
2	38.2%	81.7%	90.0%
4	16.1%	68.2%	90.2%
6	7.1%	58.2%	90.3%
8	3.5%	49.8%	90.9%

branches nested; (*AW-3*) selection cannot inflate miscoverage, now proved via the nested-band lemma and union bound of Theorem 5 rather than any conditional-coverage claim. The excess-width oracle inequality (*AW-4*) and the $O(g/K)$ modulation self-inclusion deficit—one open, one quantified—are stated in Appendix A.

6 Simulation and a design-preserving stress test

Simulation establishes two facts: that attaching the band to the wrong unit collapses coverage (Section 6.1), and that the safe selector meets its guarantees on data it has never seen. For the latter we develop the selector on one simulation grid and then confirm the *frozen* rule once on an independent design. The separation is deliberate: the development grid may set the gate constants, but it may not be reused as validation performance.

6.1 The wrong unit collapses coverage

The most basic claim deserves a direct test: a band attached to the wrong unit covers the object of interest far below nominal. On a ground-truth simulation ($K = 30$ exchangeable countries, $T = 6$ thresholds, error tensors correlated across rounds and thresholds; Table 1), we ask each band for *trajectory* coverage—the whole country curve-path covered at once. All three bands are unstudentized and finite-sample constructed, so the only thing that varies is the unit of the score.

The per-round band, marginally valid at the round level, covers the whole trajectory only 0.82 at $L = 2$ and 0.50 at $L = 8$; the marginal band collapses faster still. The country-trajectory band holds ≈ 0.90 regardless of L , because the whole trajectory is a single exchangeable unit and its sup-score obeys the finite-sample order-statistic guarantee. This is the coverage face of the paper’s thesis: the same mechanism drives the certification-count collapse in Section 8, where asking for the honest object—not asking it more strictly—removes the over-detection. Pre-registered and reported as produced (`e28_wrong_unit_coverage`).

The development grid that motivates and calibrates the safety gate, and a controlled Gaussian illustration that isolates the finite- K over-shrinkage mechanism behind the $K=320$ -only activation seen below, are given in the Supplementary Material.

Independent confirmatory validation. We froze the selector and its constants (Appendix B), sealed a preregistration and a seed manifest, and ran the deployed procedure *once* on a grid disjoint from development in K ($\{25, 40, 80, 160, 320\}$), in ρ (dense at the ρ_0 cutoff), and in data-generating process (ten families spanning skewed and heavy-tailed noise, heterogeneous design variance, weighting, irregular trajectory length, a *misspecified* noise law, and weak/strong cross-round dependence). The runner verifies the config hash against the sealed manifest before executing. The numbers are reported exactly as produced (Table 2, Figure 1).

The design-aware layer behaves as designed. In the sealed run all twelve cells below the 0.88 floor routed to base clustered PCB—the deconvolution branch never once rode into a floor violation. Under the corrected scorer the deployed pipeline passes every criterion outright: the selector reduces to PCB at low ρ (width ratio 1.005), activates deconvolution only at $K=320$ (never for $K \leq 160$; 24 cells, up to 99% of draws at the highest ρ), and there delivers a band $0.753\times$ the conservative width—wider than the sealed run’s 0.577 because the α_{dec} budget of Theorem 5 prices the branch’s estimated-law remainder honestly—with worst $K=320$ cell coverage 0.900.




Figure 1: Independent confirmatory validation on the unseen grid (10 new DGP families, new K , new ρ). Left: deployed coverage vs. $\hat{\rho}$ per K (line = mean over families, shade = worst family) against the 0.88 floor and 0.90 nominal. Right: deconvolution activates only at $K=320$; small K abstains.

The holdout exposed one base-method deficit, which we then resolved. The studentized base band (in-sample pooled modulation) undercovered at small K and low ρ under the hardest DGPs, because dividing each calibration score by a modulation estimated from the same clusters—which the unobserved target does not share—breaks the exchangeability the conformal guarantee needs. We therefore deploy the *unstudentized* clustered band as the base branch: with no data-dependent denominator the scores $R_c = \max_t |E_c(t)|$ are exchangeable, and the order-statistic band is finite-sample exact at *any* K , $\Pr(\text{cover}) \geq [(1 - \alpha)(K+1)]/(K+1) \geq 1 - \alpha$, distribution-free (Theorem 3). A preregistered rerun on a fresh grid confirms the fix at a width cost of 1.02–1.05 $\times$ on the real ESS transport bands; the granular cell-by-cell results are in the Supplementary Material. The band the paper deploys is thus finite-sample valid at the sample sizes comparative politics actually has. A second, extreme-tail limitation of the Gaussian-calibrated remainder bound $\hat{\delta}_{\text{UCB}}$ —confined to 0.02% of draws and absorbed by the conservative branch—is documented in the Supplementary Material. We do not claim uniform nominal coverage; we claim—and the holdout confirms—that the design-aware correction is never the *cause* of undercoverage and is materially more efficient than conservative where it safely activates.

Table 2: Holdout validation of the safe selector (450 cells, 800 reps each). The *sealed* run (upper rows) scored the studentized S2 base band, whose finite- K self-inclusion deficit it exposed — the finding that motivated deploying the *unstudentized* band instead. The *corrected-scorer rerun* (lower rows) scores the band the package actually returns (U0 anchors with the α -budgeted deconvolution branch of Theorem 5) on the same 450-cell grid with a fresh seed, disclosed because the originally sealed configuration file was not preserved in the [archive](#).

criterion	target	result	status
<i>Sealed run, S2 scorer (historical):</i>			
P3: design-aware never causes a sub-floor cell	sub-floor $\Rightarrow$ not deconv	all 12 sub-0.88 cells are PCB	pass
E2: deconvolution narrower than conservative	$\leq 0.90\times$	0.577	pass
P1: cells floor-compatible (95% MC)	all 450	442/450	diagnostic [†]
P2: worst-cell coverage	≥ 0.86	0.843	diagnostic [†]
<i>Corrected-scorer rerun, deployed U0/α-budget band:</i>			
P1: cells floor-compatible (95% MC)	all 450	450/450	pass
P2: worst-cell coverage	≥ 0.86	0.882	pass
E1: low- ρ width vs PCB	$\leq 1.05\times$	1.005	pass
E2: deconvolution narrower than conservative	$\leq 0.90\times$	0.753	pass
E3: deconvolution abstains at small K	rises with K	only $K=320$	pass

[†] Diagnostic of the studentized base band, not a deployed-band failure; these cells motivated the fix. The deployed unstudentized band is finite-sample exact at any K (Theorem 3); under the corrected scorer its worst holdout cell is 0.882 (1.65 MC-SE below the 0.90 nominal, i.e. binomial-noise compatible with exactness) and no cell fails the preregistered floor criterion.

Design-preserving stress test. A design-preserving subsampling of the real AmericasBarometer confirms the deconvolution regime is unreachable even semi-synthetically at survey scale, with ρ saturating near 0.23 (Supplementary Material)—direct evidence for the regime characterization of Section 7.

7 Evidence from the ESS and the AmericasBarometer

The regime characterization. On real cross-national data the honest band reduces to clustered conformal—and the reason is instructive. Across *both* surveys, *both* estimands (attitude levels and wave-to-wave change), and three outcomes (trust in parliament, satisfaction with democracy, system support), national cross-national inference is a uniformly *low design-noise* regime: the estimated $\hat{\rho} \leq 0.20$, comfortably below the $\rho_0 = 0.47$ cutoff, because between-country variation dwarfs within-survey design noise (Figure 2). Consequently the safe selector reduces to clustered PCB on *all* real data—not by assumption but by its




Figure 2: Estimated design-to-transport SD ratio $\hat{\rho}$ for the AmericasBarometer, levels and wave-to-wave change. Every country-outcome sits well below the $\rho_0 = 0.47$ cutoff: national cross-national attitude inference is a uniformly low design-noise regime, so the safe selector reduces to clustered PCB.

own low- ρ gate. The elaborate design-aware machinery is, on real cross-national attitude data, correctly inactive.

This is a substantive finding: the impossibility of Section 3 bites in principle, but the deconvolution it necessitates is never triggered at survey scale, so honest inference here *is* clustered conformal on the observed scores.

Why the deconvolution is unreachable, precisely. The inactivity is not an accident of one dataset; it is forced by two structural barriers, and we can say how far each real survey is from clearing them.

Proposition 1 (Survey-scale unreachability). *The deployed selector rides the deconvolution branch only if both a need gate $\hat{\rho}_{\text{LCB}} > \rho_0$ and a reliability gate $D \leq \tau_D$ hold, where $D = \max_t \widehat{\text{SE}}(\hat{s}_T^2)/\hat{s}_T^2$. Two distribution-free facts bound these. (a) $\rho = \sqrt{v^2}/s_{\text{plug}} \in [0, 1)$ with $s_{\text{plug}}^2 = s_R^2 + \overline{v^2}$, so the need gate can fire only when the design noise is an appreciable fraction of the between-population signal s_R . (b) Since $\hat{s}_T \leq s_{\text{plug}}$, the reliability obeys $D \geq \sqrt{2/(K-1)}$; hence the reliability gate requires $K \geq 1 + 2/\tau_D^2$, i.e. $K \geq 94$ at the frozen $\tau_D = 0.147$, exchangeable populations.*




Figure 3: The two barriers, measured on real data. *Left*: across ESS age-band subpopulations $\hat{\rho}$ saturates far below $\rho_0 = 0.47$. *Right*: the finite- K reliability D sits on its algorithmic floor $\sqrt{2/(K-1)}$ (deterministic for the frozen diagnostic; Proposition 1a); passing the reliability gate needs $K \geq 94$, which cross-national surveys ($K \leq 33$ ESS, ≤ 105 WVS) reach only where—as on WVS—the design noise is too small for the need gate.

The two requirements pull apart on real surveys (Figure 3). On the ESS, narrowing to the noisiest subpopulations (youth age bands) raises $\hat{\rho}$ from 0.08 to at most 0.29—still far below ρ_0 —while $K \leq 33$ leaves $D \geq 0.25$, so *neither* gate opens (a scan of 42 subpopulation cells activates deconvolution in none). The World Values Survey is the mirror image: with $K \approx 100$ countries the reliability gate is *feasible* for it (the floor $\sqrt{2/(K-1)}$ sits below $\tau_D = 0.147$; four of the five items measure $D = 0.139$ – 0.142 and pass, while the importance-of-democracy item sits at $D = 0.1473$, a hair’s breadth over the frozen threshold), but it ships no cluster identifiers, so its weights-only design noise gives $\hat{\rho}_{LCB} \leq 0.09$ and the need gate fails decisively. The survey large enough to *estimate* the deconvolution has almost no design noise to remove; the survey with real clustered design noise has far too few countries to estimate it. The two barriers are anti-correlated, and neither of the two largest repeated cross-national surveys clears both.

The barriers are not survey artifacts. Both gates are distribution-free, so the boundary transfers to any conformal procedure calibrated on estimated objects (Remark 1); a non-

survey MRP-style simulation confirms the reliability floor and identifies the one many-area setting, with appreciable posterior uncertainty, where the *need* gate can fire (Supplementary Material).

The AmericasBarometer supplies the design effect the ESS lacks. Its full stratified-PSU structure lets us compute a proper design bootstrap, and the design effect is real (median ≈ 1.13 – 1.20 , up to 1.9 in the most clustered country-years). Yet certification is coarse-robust: the design effect moves band *width* but rarely flips the binary decline decision, because the decision is dominated by the between-country transport spread, not the within-country design noise—the same low- ρ fact, seen from the certification side (details in the Supplementary Material).

What real data validates, and what it cannot. Real data validates the reduction to PCB, the target-blind selection, the efficiency of PCB over the conservative envelope, the real design effect, and the political reclassification of Section 8. It *cannot* validate the high- ρ deconvolution branch or the finite- K remainder—those regimes do not occur at survey scale and rest on the simulation of Section 6 and the theorems, not on the ESS and AmericasBarometer data.

8 Political reanalysis: how many democracies really declined?

The method’s payoff is a sharp distinction between what marginal analysis appears to show and what a design-aware, simultaneous band certifies. We analyze trust in parliament (`trstprl`) and satisfaction with democracy (`stfdem`) in the ESS over 2018–2022, treating each country’s attitude distribution as the object and asking, at each rung of a guarantee hierarchy, in how many countries a decline survives.

The comparative literature has itself grown skeptical of a continent-wide trust decline: where Foa and Mounk (2016) saw “democratic deconsolidation,” Norris (2017) saw trendless fluctuation, Voeten (2017) saw specification-fragile artifacts, and Wuttke et al. (2022)—analyzing eighteen democracies under preregistration—found endorsement of self-government largely intact; the most comprehensive recent accounting finds no universal secular decline in institutional trust at all (Valgarðsson et al., 2025). This qualitative consensus holds that “trust is falling everywhere” is an over-reading; our contribution is a formal instrument that says so, and the wrong-unit diagnosis of *why* it arises. A mean can fall because a mode hollows out or a tail thickens, with different substantive meanings (DiMaggio et al., 1996); requiring the whole distribution to shift is the stronger, and more honest, claim.

A hierarchy of claims. The claims differ in their unit and are not interchangeable. A *pairwise* claim asks whether some pair of waves shows a distributional shift; an *any-pair* claim aggregates these within a country; a *net* claim asks whether the first-to-last span declined; a *persistent country-wide* claim asks whether the *whole* distribution declined *simultaneously* across the trajectory—a single band that must hold over all wave pairs and all thresholds at once; and a *Bonferroni-across-countries* claim further controls the number of countries examined. Each rung is strictly stronger than the last, and conflating them is exactly the over-detection this paper is about.

The counts collapse up the hierarchy. For `trstpr1`, a marginal pairwise reading flags on the order of twenty countries; requiring an any-pair design-aware band drops this to twelve; the net first-to-last decline holds in roughly six; and the *persistent, country-wide, design-aware simultaneous* band—the honest object—is certified in **one** country, Greece, which also survives the Bonferroni correction across countries. Greece is precisely the case the substantive literature independently regards as a genuine crisis-era collapse of legitimacy (Armingeon and Guthmann, 2014; Torcal, 2014), so certifying only Greece is a credibility feature, not an anomaly. For `stfdem` the same pattern holds (roughly twenty-three marginally,




Figure 4: The count of certified declining countries collapses up the guarantee hierarchy (pairwise $\rightarrow$ any-pair $\rightarrow$ net $\rightarrow$ persistent country-wide design-aware), from about twenty marginally to one (Greece).

fourteen any-pair, and one persistent, again Greece), though there the Greek decline does *not* survive the additional Bonferroni rung—the across-countries correction is passed for trust in parliament only. The headline is that a *persistent distribution-wide* decline—the strong claim casual readings imply—is certifiable in a single country over this window (Figure 4).

Which countries are reclassified, and the honest word for it. The reclassified countries (Figure 5) are the weak-signal, higher-noise cases, where the apparent trend lies within the design uncertainty the survey warrants. We describe them as *certified by the marginal plug-in but inconclusive under the design-aware simultaneous band*—never as “false positives,” a term we reserve for the simulation where ground truth is known. Because under-coverage would make *more* countries certify, the “only Greece” finding is conservative in the direction that matters, and the deployed base band is finite-sample exact at the $K \approx 30$ of these data (Theorem 3), so the count is not an artifact of small- K under- or over-coverage.

Robustness across age groups. A preregistered age-group analysis (Appendix B) confirms the “only Greece” result is not an artifact of pooling ages—it holds within the older (50+) and middle (30–49) groups—while youth (18–29) certify zero persistent declines as




Figure 5: Countries certified by the marginal plug-in but inconclusive under the design-aware simultaneous band are the weak-signal, higher-noise cases; the apparent trend lies within the design uncertainty the survey warrants.

an underpowered null rather than youth stability, and emphatically not a story of youth-led democratic disaffection. The full age-band results are in the Supplementary Material.

A second, global target: Foa–Mounk deconsolidation. To show the correction is not a European artifact, we preregistered a reanalysis of the Foa–Mounk “democratic deconsolidation” thesis (Foa and Mounk, 2016, 2017) on the World Values Survey / EVS trends (1981–2022, 59–77 countries with ≥ 2 waves; per the replication package’s preregistration), recoding the deconsolidation battery (importance of democracy, rejection of strong-leader and army rule, support for a democratic system, confidence in parliament) so that deconsolidation is a persistent, distribution-wide decline. The verdict is calibration, not denial. A marginal wave-by-wave reading flags 36–64 countries per item; the persistent, weights-aware, distribution-wide object is certified in 6–17 (8–24% of countries)—an over-count of 2.6–6.5 $\times$. The design-aware layer demotes a further 17–40% per item ($\approx 30\%$ on average) at the persistent bar. Deconsolidation is thus *real in a specific minority* (the certified importance-of-democracy set includes Turkey, the Philippines, Mexico, and Tunisia) but is over-stated several-fold by the marginal evidence the thesis rests on. The youth “embrace authoritarianism” claim is only half-supported (Supplementary Material). WVS carries no PSU/stratum, so its survey-design correction is a weights-only respondent bootstrap; the

point is the cross-family robustness of the wrong-unit correction, from the dramatic ($20 \rightarrow 1$ for European trust) to the moderate ($2.6\text{--}6.5\times$ globally).

The certified core: where the surviving deconsolidation lives. The instrument does not only subtract; aggregating the per-item certifications yields a positive characterization of where persistent deconsolidation is concentrated (Figure 6; replication experiment E30). Of the 38 countries certified on at least one battery item, 13 form a *certified core* certified on two or more of the five: Albania on four; Bosnia and Herzegovina, Ecuador, and Tunisia on three; and Azerbaijan, Finland, Ghana, Iraq, Lebanon, Rwanda, Switzerland, Trinidad and Tobago, and Uzbekistan on two. Three features of this set reframe the debate. *First, the core is not the consolidated West.* It is dominated by post-communist states of the Balkans, Caucasus, and Central Asia and by the Arab-Spring aftermath (Tunisia, Lebanon, Iraq), with fragile Latin American and African cases alongside—young or partial democracies and post-authoritarian societies, not the mature democracies the deconsolidation thesis (Foa and Mounk, 2016, 2017) was formulated about. Consolidated Western democracies enter the core exactly twice—Finland and Switzerland—and only on the strongman/army-rule pair. *Second, the Western syndrome that does survive is specific.* Where a consolidated democracy certifies at all, it is either rising openness to authoritarian alternatives (the Finnish–Swiss pair) or falling confidence in parliament (Norway, the United Kingdom)—never a persistent distribution-wide decline in support for democracy as a system, consistent with the preregistered finding that endorsement of self-government remains intact (Wuttke et al., 2022). *Third, the celebrated backsliding cases are narrower than their reputation.* Turkey, the Philippines, and Mexico certify only on the importance-of-democracy item: their publics persistently downgrade democracy’s importance while the remaining items do not clear the persistent bar—executive aggrandizement with a partially demobilized public, rather than wholesale attitudinal collapse. The counts are descriptive aggregation of per-item certifications, each of which carries its own finite-sample guarantee; item coverage varies by country

($K = 59\text{--}77$ per item), so multi-item certification is harder where fewer items are measured, and the regional labels are descriptive groupings, not covariates.

9 Discussion

Uncertainty in claims about repeated cross-national surveys is usually computed for the wrong unit—twice: a wave-pair mean contrast stands in for a persistent distribution-wide trajectory, and an estimated distribution for the latent one it samples. We give a multiplicity-honest simultaneous band for the first and a hard boundary for the second: without the design-noise law the latent uncertainty is non-identified (Theorem 1), and *with* it the correction is still unreachable at cross-national scale, blocked by two anti-correlated barriers no survey clears at once (Proposition 1). The honest deployed object is therefore the reduction—a clustered band exact at any K (Theorem 3) inside a selector that deconvolves only when safe and otherwise abstains, provably correctly (Theorem 5), confirmed on an independent pre-registered design. The substantive payoff is carried by the two reanalyses; the design-aware layer is a diagnostic that certifies, rather than corrects, at survey scale.

When it reduces to ordinary conformal. On real cross-national attitude data the design noise is small relative to between-country spread, so the method reduces *exactly* to clustered population conformal (Section 7): the machinery is diagnostic, *certifying* that the simple band is the honest one and reporting the ρ that justifies it. The deconvolution regime requires *both* many exchangeable units ($K \geq 94$) and design noise comparable to the between-unit signal—a combination cross-national surveys never provide (Proposition 1) but many-unit clustered settings (subnational small areas, schools, clinics) can. A simulation of exactly that setting (Supplementary Material, E31) confirms the payoff is real, not hypothetical: at $K \approx 220\text{--}300$ small areas with posterior noise around 0.7 of the between-area signal, the frozen pipeline activates deconvolution on up to 98% of draws and delivers a band $0.69\text{--}0.74\times$ the conservative envelope’s width at valid coverage—the positive-regime frontier the

survey-scale impossibility leaves open.

Limitations. Without design metadata—as in the World Values Survey, which releases weights but limited cluster identifiers—only a weights-only replication is possible; the impossibility result says no design-valid inference can be recovered from the point estimates alone. And the finite- K remainder bound for the deconvolution branch was calibrated on Gaussian data, so it is conservative but not fully robust under strong dependence, pointing to a richer calibration family or a dependence-aware gate—though, since that branch is provably never reached at cross-national scale (Proposition 1), this bears only on hypothetical large- K , high-noise applications. (The base band’s earlier small- K undercoverage is *resolved*: the deployed unstudentized band is finite-sample exact at any K (Theorem 3) at $\leq 5\%$ width on real data.)

The general lesson. Attaching uncertainty to the right unit changes substantive conclusions. For European trust a marginal reading flags about twenty declining democracies; the persistent, distribution-wide, design-aware object certifies one (Greece). For the Foa–Mounk “deconsolidation” thesis the marginal evidence over-counts the persistent phenomenon by $2.6\text{--}6.5\times$ —real in a specific minority, not the broad erosion claimed, and that minority is itself concentrated in post-communist and Arab-Spring-aftermath states rather than the consolidated West (the certified core of Section 8). The pattern is not confined to surveys: wherever a conformal or prediction procedure is calibrated on estimated objects—small-area model outputs, learned features, meta-analytic effects (Remark 1)—the calibration targets carry their own error, and the claim unit is often coarser than the trend asserted. The right response is not to widen every band reflexively but to state the strongest object the data support, correct the estimation noise when it can be identified, and abstain honestly when it cannot.

A Proofs and efficiency results

Full statements and proofs of Theorems 1–5 and Proposition 1, the efficiency results (AW-1–AW-4), and the modulation self-inclusion deficit are collected in Section S1 of the online supplementary material. Each theorem is additionally paired with a machine-checked contract test that verifies its guarantee numerically.

B Frozen selector specification

The frozen constants ($\rho_0 = 0.47$, $\widehat{\delta}_{\text{UCB}}(D) = 0.0061 + 0.0943 D$, and the remainder), their derivation on the development grid, and the sealed independent-validation protocol (evaluation grid, seed manifest, holdout results) are provided in the replication package.

Acknowledgments

The author used Claude (Anthropic; Opus 4.7–4.8 and Fable 5) between 18 January and 17 July 2026 to assist with analysis and plotting code, figure generation and formatting, and the revision of proof arguments and theorem statements (including an audit that corrected an invalid selection argument and a mis-scored validation run, both disclosed in the text). The research design, methodology, and core analysis code are the author’s own. All AI-assisted output was reviewed and verified by the author, who takes full responsibility for the content of this article.

Data and Code Availability

All simulation, theory, and benchmark results are deterministic, reproducing exactly under fixed seeds with no external data, and each theorem is paired with a machine-checked contract test. The ESS, World Values Survey, and AmericasBarometer/LAPOP microdata are

licensed to their providers and not redistributed; the replication package documents the exact variables, waves, and retrieval steps. The full package—the `pcb` module, all simulations and contract tests, and the analysis pipeline—will be deposited in the Political Analysis Dataverse (Harvard Dataverse) upon conditional acceptance.

References

- Armington, K. and K. Guthmann (2014). Democracy in crisis? the declining support for national democracy in european countries, 2007–2011. *European Journal of Political Research* 53(3), 423–442.
- Barber, R. F., E. J. Candès, A. Ramdas, and R. J. Tibshirani (2023). Conformal prediction beyond exchangeability. *The Annals of Statistics* 51(2), 816–845.
- Benjamini, Y. and Y. Hochberg (1995). Controlling the false discovery rate: a practical and powerful approach to multiple testing. *Journal of the Royal Statistical Society Series B* 57(1), 289–300.
- Binder, D. A. (1983). On the variances of asymptotically normal estimators from complex surveys. *International Statistical Review* 51(3), 279–292.
- Cohen, M. J., N. Lupu, and E. J. Zechmeister (2021). The political culture of democracy in the americas, 2021: Taking the pulse of democracy. Technical report, LAPOP, Vanderbilt University. AmericasBarometer.
- DiMaggio, P., J. Evans, and B. Bryson (1996). Have americans’ social attitudes become more polarized? *American Journal of Sociology* 102(3), 690–755.
- Diquigiovanni, J., M. Fontana, and S. Vantini (2021). Conformal prediction bands for multivariate functional data. *arXiv preprint arXiv:2106.01792*.

- Dunn, R., L. Wasserman, and A. Ramdas (2023). Distribution-free prediction sets for two-layer hierarchical models. *Journal of the American Statistical Association* 118(544), 2491–2502.
- Einbinder, B.-S., S. Bates, A. N. Angelopoulos, A. Gendler, and Y. Romano (2024). Label-noise robustness of conformal prediction. *Journal of Machine Learning Research* 25. arXiv:2209.14295.
- Fan, J. (1991). On the optimal rates of convergence for nonparametric deconvolution problems. *The Annals of Statistics* 19(3), 1257–1272.
- Foa, R. S. and Y. Mounk (2016). The democratic disconnect. *Journal of Democracy* 27(3), 5–17.
- Foa, R. S. and Y. Mounk (2017). The signs of deconsolidation. *Journal of Democracy* 28(1), 5–15.
- Gelman, A., J. Hill, and M. Yajima (2012). Why we (usually) don’t have to worry about multiple comparisons. *Journal of Research on Educational Effectiveness* 5(2), 189–211.
- Gibbs, I. and E. Candès (2021). Adaptive conformal inference under distribution shift. In *Advances in Neural Information Processing Systems*, Volume 34.
- Gibbs, I., J. J. Cheng, and E. Candès (2024). Conformal prediction with conditional guarantees. *Journal of the Royal Statistical Society Series B*. arXiv:2305.12616.
- Kaminska, O. and P. Lynn (2017). Survey-based cross-country comparisons where countries vary in sample design: issues and solutions. *Journal of Official Statistics* 33(1), 123–136.
- Lee, Y., R. F. Barber, and R. Willett (2023). Distribution-free inference with hierarchical data. *arXiv preprint arXiv:2306.06342*. Forthcoming, ACM/IMS Journal of Data Science.

- Lei, J., M. G'Sell, A. Rinaldo, R. J. Tibshirani, and L. Wasserman (2018). Distribution-free predictive inference for regression. *Journal of the American Statistical Association* 113(523), 1094–1111.
- Lumley, T. (2010). *Complex Surveys: A Guide to Analysis Using R*. Wiley.
- Norris, P. (2017). Is western democracy backsliding? diagnosing the risks. HKS Faculty Research Working Paper RWP17-012, Harvard Kennedy School; Journal of Democracy web exchange. SSRN 2933655.
- Randahl, D., J. P. Williams, and H. Hegre (2026). Bin-conditional conformal prediction of fatalities from armed conflict. *Political Analysis* 34(1), 96–108.
- Sesia, M., Y. X. R. Wang, and X. Tong (2025). Adaptive conformal classification with noisy labels. *Journal of the Royal Statistical Society Series B* 87(3), 796–815.
- Tibshirani, R. J., R. F. Barber, E. Candès, and A. Ramdas (2019). Conformal prediction under covariate shift. In *Advances in Neural Information Processing Systems*, Volume 32.
- Torcal, M. (2014). The decline of political trust in spain and portugal: economic performance or political responsiveness? *American Behavioral Scientist* 58(12), 1542–1567.
- Valgarðsson, V., W. Jennings, G. Stoker, H. Bunting, D. Devine, L. McKay, and A. Klassen (2025). A crisis of political trust? global trends in institutional trust from 1958 to 2019. *British Journal of Political Science* 55. Article e15.
- Voeten, E. (2017). Are people really turning away from democracy? Journal of Democracy web exchange. SSRN 2882878.
- Vovk, V., A. Gammerman, and G. Shafer (2005). *Algorithmic Learning in a Random World*. Springer.
- Wieczorek, J. (2023). Design-based conformal prediction. *Survey Methodology* 49(2). Statistics Canada, Catalogue No. 12-001-X. arXiv:2303.01422.

Wuttke, A., K. Gavras, and H. Schoen (2022). Have europeans grown tired of democracy? new evidence from eighteen consolidated democracies, 1981–2018. *British Journal of Political Science* 52(1), 416–428.




Figure 6: Co-certification across the five deconsolidation battery items (WVS/EVS 1981–2022). Filled cells are country-items where the persistent, distribution-wide, weights-aware decline is certified; the 13 countries above the dashed line are certified on ≥ 2 items. The certified core is concentrated in post-communist and Arab-Spring-aftermath states; the consolidated West appears only via the strongman/army pair (Finland, Switzerland) and parliament confidence (Norway, United Kingdom).