

# Supplementary Material for “The Wrong Unit of Uncertainty: Simultaneous Conformal Bands for Repeated Cross-National Surveys”

The Authors

This online supplement collects secondary material referenced from the main manuscript. All content, numbers, and figures are reproduced unchanged; only figure numbers carry an “S” prefix.

## S1 Proofs

This section gives full statements and proofs of the results of the main text, in the main-manuscript numbering. Each result is paired with a machine-checked contract test in the replication package that verifies its guarantee numerically. Where the source argument is a sketch or rests on an estimated constant, we say so.

**Common setup.** We observe  $K$  exchangeable source countries  $c = 1, \dots, K$  plus a target  $K+1$ . Each country  $c$  has a latent attitude curve  $F_c$  (a CDF over a grid  $t_1 < \dots < t_T$ ) and a transport center  $\mu$  (a leave-one-out mean or fitted center), with plug-in error curve  $E_c(t) = \hat{\theta}_t(P_c) - \theta_t(P_c)$ . The latent transport score is  $R_c = \max_{r,t} |F_c(t) - \mu(t)|/\sigma(t)$ , the country-level (exchangeable-unit) score.  $F_c$  is not observed; a complex survey returns  $\tilde{F}_c = F_c + S_c$

with design error  $S_c$ , inducing the contaminated score  $\tilde{R}_c = R_c + \xi_c$  with  $\mathbb{E}[\xi_c \mid R_c] = 0$  and  $\text{Var}(\xi_c \mid \cdot) = v_c^2$ , where  $v_c^2$  is estimated by a stratified-PSU / replicate-weight design bootstrap. Write  $s_{\text{plug}}^2 = s_R^2 + \overline{v^2}$  and the design-to-total SD ratio  $\rho = \sqrt{\overline{v^2}}/s_{\text{plug}} \in [0, 1)$ . Two targets are kept distinct throughout: the survey-estimate score  $\tilde{R}_{K+1}$  (T1) and the latent score  $R_{K+1}$  (T2), the object the political claim concerns.

### S1.1 Theorem 1 (non-identification and necessity of design information)

*Statement.* Let the observed score be  $\tilde{R} = R + \xi$  with  $R \perp \xi$  and  $\xi$  mean-zero. (i) The law of  $R$ —hence its  $(1 - \alpha)$  quantile—is not identified from the law of  $\tilde{R}$  without restrictions on the law of  $\xi$ . (ii) Among bands measurable with respect to the observed scores only, none is valid for the latent target at level  $1 - \alpha$  for all admissible design laws while being strictly narrower than the plug-in band whose radius is the  $(1 - \alpha)$  conformal quantile of  $|\tilde{R}|$ .

*Proof.* (i) Convolution is not invertible without knowing the noise law. Take  $\mathcal{L}(\tilde{R}) = N(0, \sigma^2)$ . Then  $(R \sim N(0, \sigma^2), \xi \equiv 0)$  and, for any  $0 < \tau^2 < \sigma^2$ ,  $(R \sim N(0, \sigma^2 - \tau^2), \xi \sim N(0, \tau^2))$  both have  $\xi \perp R$ ,  $\xi$  mean-zero, and induce the same  $\mathcal{L}(\tilde{R})$ ; their latent  $(1 - \alpha)$  quantiles  $z_{1-\alpha}\sigma$  and  $z_{1-\alpha}\sqrt{\sigma^2 - \tau^2}$  differ. The same construction applies to any  $\mathcal{L}(\tilde{R})$  that is a nontrivial convolution. (ii) Under the member  $\xi \equiv 0$  the latent target equals  $\tilde{R}_{K+1}$ , which is exchangeable with  $\{\tilde{R}_c\}_{c \leq K}$ ; the conformal rank identity gives  $\Pr(\tilde{R}_{K+1} \leq \hat{q}) = [(1 - \alpha)(K+1)]/(K+1)$  at  $\hat{q}$  the corresponding order statistic, so any radius below the  $[(1 - \alpha)(K+1)]$ -th order statistic of  $\{|\tilde{R}_c|\}$  has latent coverage  $< 1 - \alpha$  on this member. A band narrower than the plug-in on a positive-measure event therefore undercovers the latent target on the  $\xi \equiv 0$  law it cannot distinguish from the contaminated one. In the degenerate case  $\xi \equiv 0$  the plug-in is itself the tightest valid band, so no uniform improvement exists. Thus strict efficiency is attainable only by shrinking the class of admissible laws—i.e. by supplying the law of  $\xi$  (equivalently the  $v_c$ ), which is exactly the

design-bootstrap information. □

## S1.2 Theorem 2 (oracle validity)

*Statement.* With the countries exchangeable and the design law known: (a) the clustered band on  $\{\tilde{R}_c\}$  covers the survey-estimate target T1 exactly,  $\Pr(\tilde{R}_{K+1} \leq \hat{q}) \geq 1 - \alpha$ , using exchangeability only. (b) For the latent target T2 the same band satisfies, for every  $u > 0$ ,

$$\Pr(R_{K+1} \leq \hat{q}) \geq 1 - \alpha - \Pr(\xi_{K+1} < -u) - \Pr(\tilde{R}_{K+1} \in (\hat{q} - u, \hat{q}]),$$

so the T1→T2 coverage gap is controlled by a noise-tail term and a local anti-concentration term, both of which vanish with the design-noise scale (both are  $O(\rho)$  under a bounded density of  $\tilde{R}$  near its upper quantile). Equivalently, the *noise-inflated* band with radius  $\hat{q} + u$  covers the latent target at level  $1 - \alpha - \Pr(|\xi_{K+1}| > u)$  with no anti-concentration condition. (c) The oracle deconvolution band with the true target scale has width a factor  $\sqrt{1 - \rho^2} + o(1)$  of the clustered band.

*Proof.* (a) By exchangeability of  $\{\tilde{R}_c\}_{c=1}^{K+1}$  the rank of  $\tilde{R}_{K+1}$  is uniform on  $\{1, \dots, K+1\}$ ; with  $\hat{q}$  the  $\lceil (1 - \alpha)(K+1) \rceil$ -th order statistic,  $\Pr(\tilde{R}_{K+1} \leq \hat{q}) = \lceil (1 - \alpha)(K+1) \rceil / (K+1) \geq 1 - \alpha$ , finite-sample and distribution-free, using no assumption on the design-noise law. (b) On  $\{R_{K+1} > \hat{q}\}$ , either  $\xi_{K+1} < -u$  or  $\tilde{R}_{K+1} = R_{K+1} + \xi_{K+1} > \hat{q} - u$ . Hence  $\Pr(R_{K+1} > \hat{q}) \leq \Pr(\xi_{K+1} < -u) + \Pr(\tilde{R}_{K+1} > \hat{q} - u)$ , and  $\Pr(\tilde{R}_{K+1} > \hat{q} - u) \leq \alpha + \Pr(\tilde{R}_{K+1} \in (\hat{q} - u, \hat{q}])$  by part (a). For the inflated band,  $\{R_{K+1} > \hat{q} + u\} \subseteq \{\tilde{R}_{K+1} > \hat{q}\} \cup \{|\xi_{K+1}| > u\}$  directly. (c) The oracle deconvolution scale satisfies  $s_T^2 = s^2(1 - \rho^2)$  by the law of total variance  $\text{Var}(\tilde{R}) = \text{Var}(R) + \mathbb{E}[v^2]$  under conditional mean-zero noise, giving the width factor  $\sqrt{1 - \rho^2} + o(1)$ . □

*Remark 1* (Symmetric noise alone is *not* sufficient for latent conservativeness). It is tempting to argue that symmetric design noise makes the observed-score band conservative for the latent target, on the grounds that a symmetric convolution is a mean-preserving spread.

That inference is false without shape conditions on the latent score law: a mean-preserving spread does not imply quantile dominance at a fixed level. Counterexample:  $R \in \{0, 10\}$  with masses  $(0.85, 0.15)$  and  $\xi = \pm 1$  with equal probability give  $Q_{0.9}(R) = 10$  but  $Q_{0.9}(R + \xi) = 9$  — the symmetric noise moves mass *below* the latent quantile — and the conformal band calibrated on the noisy scores undercovers the latent target (0.867 at  $K = 200$ ,  $\alpha = 0.10$ , in the replication package’s contract test `test_latent_target_bridge.py`). Bimodal latent score laws of this shape arise naturally when a minority of countries follows a distinct regime. This is why Theorem 2(b) is stated as an explicit two-term bound (small noise  $\Rightarrow$  small gap) rather than as unconditional conservativeness, and why on real data we report the design-noise scale ( $\rho \leq 0.23$  across all surveys examined, Section S4) that makes the gap negligible rather than assuming it away. Under additional shape conditions (e.g. a unimodal latent law with  $\hat{q}$  above its mode, or a log-concave latent density) the classical dominance argument does go through, but we do not rely on it.

### S1.3 Theorem 3 (exact finite- $K$ validity of the deployed base band)

*Statement.* Let the error curves be  $E_c(t) = \hat{\theta}_c(t) - \mu(t)$ , where the center  $\mu$  satisfies the *symmetric-construction condition*: it is (i) a fixed function independent of the  $K+1$  curves, or (ii) computed from a disjoint center fold, or (iii) computed *per unit* from that unit’s own auxiliary data by a common rule (e.g. the country’s own previous round, as in the deployed ESS certification), or (iv) a symmetric function of *all*  $K+1$  curves including the target’s (the surveyed-target validation mode). With the unstudentized cluster score  $R_c = \max_t |E_c(t)|$  and the countries exchangeable, for *every*  $K$  the band  $\hat{B} = \text{center} \pm \hat{q}$  with  $\hat{q}$  the  $\lceil (1-\alpha)(K+1) \rceil$ -th order statistic of  $\{R_c\}$  satisfies  $\Pr(F_{\text{new}} \in \hat{B}) \geq \lceil (1-\alpha)(K+1) \rceil / (K+1) \geq 1 - \alpha$ , distribution-free.

*Proof.* Under any of (i)–(iv) the map from the  $K+1$  curves to the  $K+1$  scores is symmetric in the units (conditional on the fold in case (ii)), so the scores  $E_1, \dots, E_{K+1}$  are exchangeable elements of  $\mathbb{R}^T$ , and by continuity the scores  $\{R_c\}$  are a.s. distinct. Therefore the rank of

$R_{K+1}$  is uniform on  $\{1, \dots, K+1\}$ , and with  $\hat{q} = R_{(m)}$ ,  $m = \lceil (1-\alpha)(K+1) \rceil$ ,  $\Pr(R_{K+1} \leq \hat{q}) = m/(K+1) \geq 1 - \alpha$ . The simultaneous coverage event is pointwise equivalent to the scalar event:  $F_{\text{new}}(t) \in \hat{B}(t) \forall t \iff |E_{K+1}(t)| \leq \hat{q} \forall t \iff \max_t |E_{K+1}(t)| \leq \hat{q} \iff R_{K+1} \leq \hat{q}$ . Hence simultaneous coverage is  $\geq \lceil (1-\alpha)(K+1) \rceil / (K+1)$ , distribution-free. Intersecting the band with  $[0, 1]$  and applying the isotonic tightening  $L(t) = \text{clip}_{[0,1]}(\max_{t' \leq t} \text{lo}(t'))$ ,  $U(t) = \text{clip}_{[0,1]}(\min_{t' \geq t} \text{hi}(t'))$  only shrink the band while still containing the monotone  $[0, 1]$ -valued truth, preserving validity. The split-modulation studentized variant (modulation computed on a disjoint fold) is exact by the identical argument conditional on the modulation fold.  $\square$

*Remark 2* (Transport centers estimated from the calibration pool). A grand mean over the calibration countries, used as the center for an *unsurveyed* target, violates the symmetric-construction condition: each calibration country sits inside its own center (deviations shrunk by  $(K-1)/K$ ) while the target sits outside, which biases  $\hat{q}$  downward — an anti-conservative  $O(1/K)$  asymmetry. The deployed remedy is the leave-one-out deviation  $E_c = \hat{\theta}_c - \text{mean}_{c' \neq c} \hat{\theta}_{c'}$  (`pcb.inference.conformal_band.loo_deviations`): each calibration score then excludes its own unit exactly as the target's score excludes the target. The residual asymmetry (a pool of  $K-1$  versus  $K$  units) makes calibration deviations, if anything, slightly *more* dispersed than the target's — an  $O(1/K^2)$  remainder in the conservative direction — and vanishes entirely under constructions (i)–(iv), which the headline ESS certification satisfies through its own-previous-round centers (case iii).

*Modulation self-inclusion deficit (proof sketch; estimated constant)*. If instead one studentizes by an in-sample modulation  $\hat{s}(t) = \text{SD}_c E_c(t)$  computed from the calibration curves themselves, the unobserved target does not share this denominator, which breaks exchangeability. A first-order expansion of the order-statistic bias in the number  $g$  of modulation slots each country contributes to gives a coverage deficit  $\approx 0.382(g/K)$ ; the constant 0.382 is estimated (empirical  $R^2 = 0.923$ ), not derived, and is flagged as such. Pooled modulation ( $g = 1$ ) is safe; slotwise ( $g = L$ ) undercovers by  $\approx 0.382 L/K$ . This  $O(g/K)$  deficit is why the deployed band is unstudentized.

## S1.4 Theorem 4 (estimated-law validity)

*Statement.* With the design law estimated from a size- $B$  stratified-PSU bootstrap and  $K$  countries, the deconvolution band satisfies  $\Pr(F_{\text{new}} \in \widehat{B}) \geq 1 - \alpha - \varepsilon_{K,B}$ , where  $\varepsilon_{K,B} = O(K^{-1/2}) + O(B^{-1/2}) \rightarrow 0$ .

*Proof sketch.* By the law of total variance  $\text{Var}(\widetilde{R}) = \text{Var}(R) + \mathbb{E}[v^2]$  under conditional mean-zero noise; subtracting a consistent estimate  $\widehat{v}^2$  identifies  $s_R^2$ , and the deconvolved scale  $\widehat{s}_T^2 = \max(s_{\text{plug}}^2 - \widehat{v}^2, \text{floor})$ —in the deployed finite- $K$ -safe form, with a *lower* confidence bound subtracted so the correction cannot over-shrink—recovers the target scale. The remainder decomposes as (i)  $O(B^{-1/2})$  Monte-Carlo error in  $\widehat{v}^2$ , (ii)  $O(K^{-1/2})$  finite- $K$  quantile and variance estimation, and (iii) a higher-order Gaussian-shape term; each  $\rightarrow 0$  and none is claimed exact. When  $\widehat{v}^2 \geq s_{\text{plug}}^2$  (design noise  $\gtrsim$  transport signal,  $\rho \gtrsim 1$ ) the scale hits the floor and the selector hands off to the conservative branch. This is a proof sketch: the constants in  $\varepsilon_{K,B}$  are order-of-magnitude, not tight.  $\square$

## S1.5 Theorem 5' (safe-adaptive finite- $K$ validity—the deployed pipeline)

The deployed validity argument uses *no conditional-on-selection step*. (An earlier draft argued that conditioning on a calibration-measurable selection event leaves each branch's marginal guarantee intact; that inference is false in general — the selection event and the conformal quantile are functions of the same calibration scores, so conditioning breaks the rank-uniformity that the marginal guarantee rests on. The architecture below removes the need for any such step.)

**Lemma 1** (Free selection among nested bands). *Let  $\widehat{B}_1 \subseteq \widehat{B}_2$  be two bands (as random sets) with  $\Pr(F_{\text{new}} \in \widehat{B}_1) \geq 1 - \alpha$ , and let  $\widehat{J} \in \{1, 2\}$  be any selection rule, measurable or adversarial, data-dependent or not. Then  $\Pr(F_{\text{new}} \in \widehat{B}_{\widehat{J}}) \geq 1 - \alpha$ .*

*Proof.*  $\{F_{\text{new}} \notin \widehat{B}_{\widehat{J}}\} \subseteq \{F_{\text{new}} \notin \widehat{B}_1\}$  pointwise: if the selected band misses, then a fortiori the narrower band  $\widehat{B}_1$  misses (on  $\{\widehat{J} = 2\}$  because  $\widehat{B}_1 \subseteq \widehat{B}_2$ , on  $\{\widehat{J} = 1\}$  trivially). Take

probabilities. □

*Statement (Theorem 5').* Deploy the selector of Algorithm 1 with the unstudentized PCB and conservative branches (scores  $\max_t |E_c(t)|$  and  $\max_t (|E_c(t)| + z \widehat{v}_c(t))$  respectively, both calibrated at level  $\alpha_{\text{anchor}}$ ) and the deconvolution branch calibrated at its own budget  $\alpha_{\text{dec}}$ , where  $(\alpha_{\text{anchor}}, \alpha_{\text{dec}}) = (\alpha, 0)$  when gate B is infeasible at the observed  $K$  (Proposition 1:  $K < 94$  at the frozen constants, which covers every repeated cross-national survey) and  $(0.9\alpha, 0.1\alpha)$  otherwise. Then, finite-sample and with no conditions on the selection rule,

$$\Pr(F_{\text{new},r}(t) \in \widehat{B}(r, t) \ \forall r, t) \geq \begin{cases} 1 - \alpha, & K < 94 \text{ (deconvolution unreachable),} \\ 1 - \alpha - \varepsilon_{K,B}, & K \geq 94, \end{cases}$$

where  $\varepsilon_{K,B}$  is the estimated-law remainder of Theorem 4 attached to the deconvolution branch only, reported at deployment through the observable diagnostic  $\widehat{\delta}_{\text{UCB}}(D)$ . For the survey-estimate target T1 the  $K < 94$  statement is exact; for the latent target T2 it carries the additive T1→T2 gap of Theorem 2(b), negligible at the observed  $\rho \leq 0.23$ .

*Proof.* The conservative score dominates the PCB score pointwise ( $|E_c(t)| + z\widehat{v}_c(t) \geq |E_c(t)|$ ,  $\widehat{v} \geq 0$ ), so at the common level  $\alpha_{\text{anchor}}$  the conservative order statistic is at least the PCB one and the two bands are nested:  $\widehat{B}_{\text{PCB}} \subseteq \widehat{B}_{\text{con}}$ . By Theorem 3 the PCB band is exact at  $\alpha_{\text{anchor}}$ ; by Lemma 1 any choice between the two nested branches retains  $\Pr(\text{miss}) \leq \alpha_{\text{anchor}}$ . When gate B is infeasible the selector can only ever choose between these two (Proposition 1 is a *deterministic* property of the frozen diagnostic, Section below), so the deployed guarantee is  $1 - \alpha_{\text{anchor}} = 1 - \alpha$  with no remainder and no selection conditions. When gate B is feasible the deconvolution branch (not nested; narrower) is charged its own miss budget through a union bound:

$$\Pr(\text{miss}) = \Pr(\text{miss} \wedge \widehat{J} \neq \text{dec}) + \Pr(\text{miss} \wedge \widehat{J} = \text{dec}) \leq \alpha_{\text{anchor}} + \Pr(\text{miss}_{\text{dec}}) \leq \alpha_{\text{anchor}} + \alpha_{\text{dec}} + \varepsilon_{K,B},$$

the first term by the nested-selection argument restricted to  $\{\widehat{J} \neq \text{dec}\}$  (the event  $\{\text{miss} \wedge \widehat{J} \neq \text{dec}\}$  implies a miss of  $\widehat{B}_{\text{PCB}}$ ), the second by Theorem 4's *marginal* guarantee for the always-deconvolve band at level  $\alpha_{\text{dec}}$ . No term conditions on  $\widehat{J}$ .  $\square$

*Remark 3.* The gates (A–D) therefore no longer carry any validity burden: they are an efficiency-and-robustness device that decides *which* of three individually budgeted constructions to return. This is weaker machinery than the earlier conditional claim and strictly stronger guarantees: at survey scale the deployed band is exact, full stop, and the only remainder anywhere in the pipeline is Theorem 4's  $\varepsilon_{K,B}$ , confined to the  $K \geq 94$  regime where deconvolution can activate and reported through  $\widehat{\delta}_{\text{UCB}}(D)$  when it does.

## S1.6 Proposition 1 (survey-scale unreachability)

*Statement.* The deployed selector rides deconvolution only if a *need* gate  $\widehat{\rho}_{\text{LCB}} > \rho_0$  and a *reliability* gate  $D \leq \tau_D$  both hold, with  $D = \max_t \widehat{\text{SE}}(\widehat{s}_T^2)/\widehat{s}_T^2$  computed by the frozen formula below. Two facts bound these. (*a, algorithmic — deterministic.*)  $\rho = \sqrt{\overline{v^2}}/s_{\text{plug}} \in [0, 1)$  with  $s_{\text{plug}}^2 = s_R^2 + \overline{v^2}$ ; and the deployed diagnostic satisfies  $D \geq \sqrt{2/(K-1)}$  *identically, for every possible dataset*, so gate B cannot open unless  $K \geq 1 + 2/\tau_D^2$ , i.e.  $K \geq 94$  at the frozen  $\tau_D = 0.147$ . This is a property of the frozen procedure, not a statistical estimate: the deployed pipeline provably never deconvolves at cross-national-survey  $K$ , which is what Theorem 5' uses. (*b, statistical — fourth-moment dependent.*) The interpretation of the floor as *the* sampling error of a  $K$ -cluster variance is exact under Gaussian scores and otherwise holds up to the kurtosis: the true relative SE of  $\widehat{s}^2$  is  $\sqrt{(\kappa - 1)/K} (1 + o(1))$  with  $\kappa$  the score kurtosis, larger than the Gaussian formula for heavy-tailed scores ( $\kappa > 3$ ) and smaller for platykurtic ones ( $\kappa < 3$ ).

*Proof.* (a)  $\rho^2 = \overline{v^2}/(s_R^2 + \overline{v^2}) \in [0, 1)$ , so the need gate can fire only when the design noise  $\overline{v^2}$  is an appreciable fraction of the between-population signal  $s_R^2$ . The frozen diagnostic computes  $\widehat{\text{SE}}(\widehat{s}_T^2) = \sqrt{2s_{\text{plug}}^4/(K-1) + (\text{SD}_K(v^2)/\sqrt{K})^2}$  and, by construction of the finite-

$K$ -safe scale,  $\widehat{s}_T \leq s_{\text{plug}}$  for every  $t$ ; hence

$$D = \max_t \frac{\widehat{\text{SE}}(\widehat{s}_T^2)}{\widehat{s}_T^2} \geq \frac{\sqrt{2s_{\text{plug}}^4/(K-1)}}{s_{\text{plug}}^2} = \sqrt{\frac{2}{K-1}}$$

as an algebraic identity of the deployed statistic (machine-checked in `test_prop1_floor.py`). Hence  $D \leq \tau_D$  forces  $\sqrt{2/(K-1)} \leq \tau_D$ , i.e.  $K \geq 1 + 2/\tau_D^2$ . With the frozen gate-B constants  $\widehat{\delta}_{\text{UCB}}(D) = a + bD$ ,  $a = 0.0061$ ,  $b = 0.0943$ ,  $\delta_{\text{max}} = 0.02$ , the threshold is  $\tau_D = (\delta_{\text{max}} - a)/b = 0.147$ , giving  $K \geq 1 + 2/0.147^2 \approx 94$ . Repeated cross-national surveys have  $K \leq 33$  (ESS) and  $\sim 20$ –40 (LAPOP, Afrobarometer, Eurobarometer), so the gate-B feasible region is empty at survey scale regardless of  $\rho$ : even if the need gate were passed, deconvolution stays blocked. (b) is the standard  $\text{Var}(\widehat{s}^2) = (\mu_4 - \sigma^4(K-3)/(K-1))/K$  expansion.  $\square$

*Remark 4.* Part (a) is what the validity theorem needs, and it is unconditional. Part (b) is what the *calibration* of  $\widehat{\delta}_{\text{UCB}}$  needs, and its fourth-moment dependence is a real limitation in one direction only: for platykurtic scores the frozen formula over-states the sampling error, so the gate is merely over-conservative (an efficiency cost); for heavy-tailed scores it under-states it, so  $\widehat{\delta}_{\text{UCB}}$  can under-budget the remainder — exactly the 0.02% extreme-tail failure mode the holdout exposed (Section S3), now confined by Theorem 5' to the  $K \geq 94$ ,  $\alpha_{\text{dec}}$ -budgeted branch. The exact threshold  $K \approx 94$  depends on the frozen gate-B constants (conservatively calibrated on the development grid); a less conservative  $\delta_{\text{max}}$  would lower it, but the calibration-free statement is the floor itself. The barrier is specific to the *cross-national* application: many-unit settings (subnational areas, schools, clinics) routinely have  $K \geq 94$ , where the reliability gate can open.

## S1.7 Efficiency results AW-1–AW-3 and the open oracle inequality (AW-4)

Write the three branch half-widths at level  $1 - \alpha$  as  $W_P = q_P s$  (PCB),  $W_D = q_D s \sqrt{1 - \rho^2}$  (deconvolution), and  $W_C = q_C s$  (conservative).

*AW-1 (low- $\rho$  reduction, proven).* As  $\rho \rightarrow 0$ ,  $\widehat{s}_T = s \sqrt{1 - \rho^2} \rightarrow s$  and the deconvolution scores converge pointwise (at fixed  $K$ ) to the PCB scores, so  $q_D/q_P \rightarrow 1$  by continuity of the finite-sample quantile in the studentization, and

$$W_D/W_P = (q_D/q_P) \sqrt{1 - \rho^2} = 1 - \frac{1}{2}\rho^2 + O(\rho^4)$$

by Taylor expansion. Below  $\rho_0$  the selector returns PCB exactly, so  $W_{\text{adaptive}} = W_P$ .

*AW-2 (conservative dominance, conditional).* For  $\rho \geq \rho_0$  with stable deconvolution,  $W_D/W_C = (q_D/q_C) \sqrt{1 - \rho^2}$ . The worst-case scores  $(|E| + zv)/s$  stochastically dominate  $|E|/s$ , so  $q_C \geq q_P$  and to first order  $q_C \approx q_P(1 + z\kappa\rho)$  with  $\kappa \in (0, 1]$  the correlation of the  $|E|$  and  $v$  orderings across the calibration set. Hence  $W_D/W_C \approx \sqrt{1 - \rho^2}/(1 + z\kappa\rho) < 1$  for all  $\rho > 0$ , strictly and decreasing in  $\rho$ : the conservative fallback is a genuine upper band, never wider than the Bonferroni envelope and at least as wide as PCB, and deconvolution strictly narrows it whenever design noise is present. (Conditional on the stated first-order expansion of  $q_C$ .)

*AW-3 (selection cannot inflate miscoverage, proven).* Superseded by Theorem 5': the nested-branch lemma (any selection between  $B_{\text{PCB}} \subseteq B_{\text{con}}$  is free) plus the union bound over the  $\alpha_{\text{dec}}$ -budgeted deconvolution branch give  $\Pr(\text{target} \in B_{\widehat{J}}) \geq 1 - \alpha - \varepsilon_{K,B} \mathbf{1}\{K \geq 94\}$  with *no* appeal to conditional-on-selection coverage. (The earlier mixture argument, which asserted that each branch retains its marginal guarantee conditional on  $\{\widehat{J} = j\}$ , was invalid: selection event and quantile share the calibration data.)

*AW-4 (excess-width oracle inequality, open).* Let  $W^* = \min(W_P, W_D, W_C)$ . Conjecture:  $W_{\text{adaptive}} \leq W^* + r_{K,B}$  with  $r_{K,B} = O(K^{-1/2}) + O(B^{-1/2})$  under Lipschitz width-in- $\rho$ , the

boundary zone where the target-blind  $\rho_0$  disagrees with the oracle crossover having width  $O(K^{-1/2})$  because  $\hat{\rho}$  is  $\sqrt{K}$ -consistent. This remains open: (i)  $\rho_0 = 0.47$  is fixed from the development grid rather than derived as the oracle crossover  $W_D = W_C$ , and (ii) the Lipschitz constant of width-in- $\rho$  is not established. It is stated as a candidate with a proof sketch, not a theorem.

## S2 The finite- $K$ mechanism and the development grid

**Development: why the third gate is needed.** On a Gaussian grid with known truth ( $K \in \{30, 60, 120, 240\}$ ,  $\rho$  up to 1.8), the selector transitions from PCB to deconvolution at  $\rho_0$  and to conservative at large  $\rho$ ; but a naive deconvolution that omits gate B *undercovers* at small  $K$ —coverage falls to 0.75 at  $K = 30$  in the transition zone—because the target scale is estimated from few countries (Figure S3). Adding the finite- $K$ -safe scale lifts this to 0.82, and the reliability gate B lifts the deployed worst cell to 0.862 while abstaining to conservative wherever the remainder cannot be bounded. These four transition cells at 0.862–0.878 are the development evidence that *motivates* the safety gate and *calibrates*  $\hat{\delta}_{\text{UCB}}$ ; they are not reused below.

**The finite- $K$  mechanism, isolated.** A controlled Gaussian illustration (design SD = transport SD,  $\rho=0.9$ ; `fig_theory_illustration.py`) makes the source of the undercoverage explicit and shows why the deployed rule is a selector rather than a deconvolver. Both the naive and the finite- $K$ -safe deconvolved score scales lie *below* the oracle  $s_R/s_{\text{plug}}$  at small  $K$  (Figure S1): subtracting a design-noise variance estimated from few countries over-shrinks the scale, and the resulting band is too narrow. The safe guard shrinks less, and both converge to the oracle as  $K \rightarrow \infty$ . The coverage consequence is severe and grows with the trajectory length  $L$  that the band must hold simultaneously (Figure S2): naive deconvolution falls to 0.43 at  $K=25$ ,  $L=16$ , and the finite- $K$ -safe scale improves it but does not by itself reach nominal at small  $K$ . The deployed adaptive selector detects that the remainder cannot




Figure S1: Score-scale shrinkage ratio  $\hat{s}/s_{\text{plug}}$  vs  $K$  (Gaussian illustration,  $\rho=0.9$ ,  $L=8$ ). Naive and finite- $K$ -safe deconvolution both over-shrink below the oracle at small  $K$ —the source of the finite- $K$  undercoverage—with the safe guard shrinking less; both converge as  $K$  grows.

be bounded and abstains to the conservative branch, remaining valid throughout at a width cost—the mechanism behind the  $K=320$ -only activation in the confirmatory holdout in the main text.

### S3 Extended robustness of the remainder bound

**Corrected-scorer holdout rerun.** A code audit found that the sealed holdout runner scored the PCB branch with the S2 studentized quantiles — a different band from the U0 construction the package deploys — so the sealed run certified a band the pipeline does not ship. We disclose this and report both: the sealed run (historical; it is what exposed the S2 deficit) and a full rerun of the same 450-cell grid scoring the band `dapcb` actually returns, under the revised  $\alpha$ -budget architecture of Theorem 5' and a fresh seed (the original sealed configuration file was not preserved in the archive; the grid is reconstructed exactly from the sealed run's cell table). The rerun passes every preregistered criterion outright: worst cell 0.882 (binomial-compatible with the exact level), 450/450 cells floor-compatible,




Figure S2: Latent-target coverage of the deconvolution branches vs  $K$ , faceted by trajectory length  $L$  (Gaussian illustration,  $\rho=0.9$ ; nominal 0.90 dotted). Naive deconvolution undercovers severely at small  $K$  and worse as  $L$  grows; the finite- $K$ -safe scale improves but does not itself reach nominal at small  $K$ . The deployed adaptive selector abstains to conservative and stays valid throughout.

deconvolution activating only at  $K=320$ , and a low- $\rho$  width ratio of 1.005 to plain PCB.

**Cell-by-cell resolution of the base-method deficit.** On the holdout the studentized band’s worst deployed cell was 0.843 (weak cross-round dependence at  $K=25$ ), a finite- $K$  self-inclusion deficit worst at small  $K$  and long trajectories. After switching to the unstudentized clustered band, a preregistered rerun on a fresh grid ( $K \in \{15, \dots, 60\}$ , ten DGPs) confirms the fix: the worst cell rises from 0.802 (S2,  $K=15$ ) to 0.893, every cell is compatible with its exact floor within the preregistered Monte-Carlo tolerance (floor minus 2 MC-SE; 29 of 120 cells sit numerically below the floor by at most 1.7 MC-SE, as expected from binomial noise around an exact level), and the width cost is 1.02–1.05 $\times$  on the real ESS transport bands (trust-CDF errors are  $\approx$ homoskedastic across the low-trust core, so the constant-width band is competitive; the only cells where it pays more than 10% are synthetic heavy-tailed curves at  $K=15$ , where the studentized alternative is itself invalid).

**Extreme-tail behavior of the remainder bound.** Beyond the base-method deficit resolved by deploying the unstudentized clustered band, the holdout exposed a second, extreme-tail limitation. The Gaussian-calibrated remainder bound  $\widehat{\delta}_{\text{UCB}}$  is not robust at




Figure S3: Development grid (Gaussian, known truth). The selector transitions PCB → deconvolution → conservative as  $\rho$  grows; naive deconvolution undercovers at small  $K$  in the transition zone—the finite- $K$  remainder the safety gate controls.

the extreme: in 77 of 360,000 draws (0.02%), on strongly-dependent or heavy-tailed DGPs at  $K=320$ , draws that pass gate B undercover to  $\approx 0.79$ . In the  $\rho = 0.90$  cells the gates route 96–98% of draws to conservative and the deployed cell coverage is 0.98–0.996; at  $\rho = 0.72$  the gates route far fewer draws away (deployed cell coverage 0.89–0.90 on the two hardest families), so the protection is concentration-dependent, not uniform over the  $\rho \geq 0.72$  range. Under the revised Theorem 5' this entire failure mode is confined to the  $K \geq 94$ ,  $\alpha_{\text{dec}}$ -budgeted branch and cannot touch the survey-scale guarantee.

## S4 Additional empirical robustness

**Design-preserving stress test.** Subsampling the real AmericasBarometer while preserving its stratified-PSU structure, the estimated  $\rho$  saturates near 0.23: since  $s_{\text{plug}}^2 = s_R^2 + \overline{v^2}$ ,

shrinking the sample inflates design noise and observed spread in lockstep. The deconvolution regime is therefore unreachable even semi-synthetically at survey scale—direct evidence for the regime characterization of the main text.

**The barriers are not survey artifacts.** Both gates are distribution-free, so the boundary should transfer to any conformal procedure calibrated on estimated objects (the general remark of the main text). A simulation confirms it outside the survey setting (`e29_beyond_surveys`): across Gaussian, heavy-tailed, skewed, and a non-survey model-based small-area (MRP-style) calibration, the reliability floor  $D \geq \sqrt{2/(K-1)}$  holds in every case (minimum  $D/\text{floor} \geq 1.06$ ), and the gate opens only at  $K \approx 130$ – $150$ —if anything a stronger barrier than the floor-implied  $K \geq 94$ . The one substantive difference is instructive: when a small area’s posterior noise is comparable to the between-area signal ( $\hat{\rho}_{\text{LCB}} = 0.54$  at a noise-to-signal ratio of 0.5), the *need* gate does fire, unlike any cross-national survey. A many-area MRP application with appreciable posterior uncertainty is therefore the setting where the deconvolution could actually activate—the positive-result frontier the survey-scale impossibility leaves open.

**The AmericasBarometer supplies the design effect the ESS lacks.** The ESS releases design weights but limited cluster identifiers; the AmericasBarometer releases full stratified-PSU structure, letting us compute a proper design bootstrap. The design effect is real: the ratio of the proper stratified-PSU standard deviation to a naive independent-resampling one has median  $\approx 1.13$ – $1.20$  and reaches 1.9 in the most clustered country-years. The proper band captures clustering variance the naive bootstrap misses. Yet certification is coarse-robust: the design effect moves band *width* but rarely flips the binary decline decision, because the decision is dominated by the between-country transport spread, not the within-country design noise—the same low- $\rho$  fact, seen from the certification side.

**Robustness across age groups (ESS).** A preregistered age-group analysis (preregistered in the replication package) reruns the same certification within fixed age bands, changing nothing else. The “only Greece” result is not an artifact of pooling ages: it holds within the older (50+) group for both outcomes and within the middle (30–49) group for trust in parliament, so the persistent decline is concentrated in the older and middle adult population rather than created by averaging. Youth (18–29) certify *zero* persistent declines, but we read this as an underpowered null rather than youth stability: only 17 of 30 countries clear the preregistered minimum-sample gate for a single youth band, and the smaller, noisier youth cells yield wider design-aware bands that certify less—the small- $K$  scope condition of the simulation appearing on real data. The analysis thus supports the headline while adding an honest caveat, and it is emphatically *not* a story of youth-led democratic disaffection.

**Youth deconsolidation on the World Values Survey.** The youth claim in the Foa–Mounk reanalysis is only half-supported: the young show *more* persistent decline in rating democracy essential (7 vs. 3 countries) but *not* in openness to strongman or army rule (0–1 vs. 4), against the “youth embrace authoritarianism” reading.